

```

%CROSS_VALIDATION Simple k-fold cross-validation for DT / RF function
meanAcc = cross_validation(X, y, modelType, params, kFolds)

if nargin < 5
    kFolds = 5;
end

% Create partition for k-fold CV
cv = cvpartition(y, 'KFold', kFolds); acc
= zeros(kFolds, 1);

for k = 1:kFolds
    % Get indices for this fold
    trainIdx = training(cv, k);
    testIdx = test(cv, k);

    % Split data according to indices
    Xtr = X(trainIdx, :);
    ytr = y(trainIdx);
    Xte = X(testIdx, :);
    yte = y(testIdx);

    switch lower(modelType)
        case 'dt'
            % Train a classification tree with specified max splits
            mdl = fitctree(Xtr, ytr, ...
                'MaxNumSplits', params.maxNumSplits);

            % Predict labels for validation fold
            yPred = predict(mdl, Xte);

        case 'rf'
            % Train a bagged ensemble for classification
            mdl = TreeBagger(params.numTrees, Xtr, ytr, ...
                'Method', 'classification', 'OOBPrediction', 'Off');

            % Same prediction but treeBagger.predict often returns cell arrays -> convert to numeric
            yPred = predict(mdl, Xte);
            if iscell(yPred)
                yPred = str2double(yPred);
            end

        otherwise
            error('Unknown modelType: %s', modelType);
    end
end

```

```

% Compute fold accuracy
acc(k) = mean(yPred == yte);
end

% Return mean accuracy across folds
meanAcc = mean(acc);
end

%-----

%METRICS Compute basic classification metrics for multi-class MNIST.
function metricsStruct = metrics(yTrue, yPred)

% Ensure column vectors
yTrue = yTrue(:);
yPred = yPred(:);

% Accuracy
accuracy = mean(yTrue == yPred);

% Confusion matrix
classes = unique(yTrue);
C = confusionmat(yTrue, yPred, 'Order', classes);

% Per-class precision/recall/F1
numClasses = numel(classes);
precision = zeros(numClasses, 1);
recall = zeros(numClasses, 1); f1
= zeros(numClasses, 1);

% Compute metrics for each class for
i = 1:numClasses
    tp = C(i, i);
    fp = sum(C(:, i)) - tp;
    fn = sum(C(i, :)) - tp;

    precision(i) = tp / max(tp + fp, eps);
    recall(i) = tp / max(tp + fn, eps);
    f1(i) = 2 * precision(i) * recall(i) / max(precision(i) + recall(i), eps);
end

% Package metrics into struct
metricsStruct = struct();
metricsStruct.accuracy = accuracy;

```

```

metricsStruct.confusion = C;
metricsStruct.classes  = classes;
metricsStruct.precision = precision;
metricsStruct.recall   = recall;
metricsStruct.f1       = f1;
end

%-----

% LOAD_DATA Reads MNIST IDX files into MATLAB matrices.
function [XTrain, yTrain, XTest, yTest] = load_data()

    % File paths
    thisDir = fileparts(mfilename('fullpath'));
    baseDir = fullfile(thisDir, '..', 'data');

    trainImagesFile = fullfile(baseDir, 'train-images.idx3-ubyte');
    trainLabelsFile = fullfile(baseDir, 'train-labels.idx1-ubyte');
    testImagesFile  = fullfile(baseDir, 't10k-images.idx3-ubyte');
    testLabelsFile  = fullfile(baseDir, 't10k-labels.idx1-ubyte');

    % Verify files exist
    assert(isfile(trainImagesFile), 'File not found: %s', trainImagesFile);
    assert(isfile(trainLabelsFile), 'File not found: %s', trainLabelsFile);
    assert(isfile(testImagesFile), 'File not found: %s', testImagesFile);
    assert(isfile(testLabelsFile), 'File not found: %s', testLabelsFile);

    % Load training data
    XTrain = load_images(trainImagesFile);
    yTrain = load_labels(trainLabelsFile);

    % Load test data
    XTest = load_images(testImagesFile);
    yTest = load_labels(testLabelsFile);

    fprintf('Loaded MNIST: %d train samples, %d test samples.\n', ...
        size(XTrain,1), size(XTest,1));
end

% Helper function to load images function
images = load_images(filename)
fid = fopen(filename,'rb');
if fid < 0, error('Could not open %s', filename); end

magic = fread(fid,1,'int32','ieee-be');

```

```
if magic ~= 2051
    error('Invalid magic number in MNIST image file %s',filename);
end
```

```
numImages = fread(fid,1,'int32',0,'ieee-be');
numRows  = fread(fid,1,'int32',0,'ieee-be');
numCols  = fread(fid,1,'int32',0,'ieee-be');
```

```
rawData = fread(fid, numImages*numRows*numCols, 'uint8');
fclose(fid);
```

```
rawData = reshape(rawData, numRows*numCols, numImages)';
images = double(rawData); % convert to double
end
```

% Helper function to load labels function

```
labels = load_labels(filename)
fid = fopen(filename,'rb');
if fid < 0, error('Could not open %s', filename); end
```

```
magic = fread(fid,1,'int32',0,'ieee-be'); if
magic ~= 2049
    error('Invalid magic number in MNIST label file %s',filename);
end
```

```
numLabels = fread(fid,1,'int32',0,'ieee-be');
labels = fread(fid, numLabels, 'uint8');
fclose(fid);
end
```

```
-----
```

%PREPROCESS_DATA Basic preprocessing for MNIST pixels function

```
[XTrainProc, XTestProc] = preprocess_data(XTrain, XTest)
```

```
fprintf('Preprocessing data (normalisation)...\\n');
```

```
XTrainProc = double(XTrain);
XTestProc = double(XTest);
```

```
% scale pixels from [0,255] to [0,1]
XTrainProc = XTrainProc / 255;
XTestProc = XTestProc / 255;
```

% visualize pixel intensity distribution

```

figure;
histogram(XTrainProc(:, 50);
xlabel('Pixel intensity (rescaled to [0,1])');
ylabel('Frequency');
title('Distribution of rescaled pixel intensities');
grid on;

% normal standard distribution
mu = mean(XTrainProc, 1);
sigma = std(XTrainProc, 0, 1) + 1e-6;
XTrainProc = (XTrainProc - mu) ./ sigma;
XTestProc = (XTestProc - mu) ./ sigma;

% visualize standardised pixel distribution
figure;
[f,xi] = ksdensity(XTrainProc(:));
plot(xi,f,'LineWidth',2);
xlabel('Pixel value (z-score)');
ylabel('Density');
title('Density of standardised pixel values');
grid on;

end

%-----



%TRAIN_DECISION_TREE Train a decision tree classifier with simple CV
function [dtModel, info] = train_decision_tree(XTrain, yTrain)

fprintf('\n==== Training Decision Tree ====\n');

% Hyperparameter grid
hyperparamValues = [100, 200, 400, 600, 800];

kFolds = 5;
cvAccuracy = zeros(size(hyperparamValues));

% Time the cross validation
tCV = tic;
for i = 1:numel(hyperparamValues)
    maxSplits = hyperparamValues(i);
    fprintf(' CV for MaxNumSplits = %d ... \n', maxSplits);

    params.maxNumSplits = maxSplits;
    cvAccuracy(i) = cross_validation(XTrain, yTrain, 'dt', params, kFolds);
end

```

```

end
cvTimeSeconds = toc(tCV);

% Choose best hyperparameter
[~, bestIdx] = max(cvAccuracy);
bestMaxSplits = hyperparamValues(bestIdx);
fprintf('Best MaxNumSplits = %d (CV accuracy = %.3f)\n', ...
    bestMaxSplits, cvAccuracy(bestIdx));

% Time the final training
tTrain = tic;
dtModel = fitctree(XTrain, yTrain, 'MaxNumSplits', bestMaxSplits);
trainTimeSeconds = toc(tTrain);

% Fill info structure
info.modelName      = 'Decision Tree';
info.hyperparamName = 'MaxNumSplits';
info.hyperparamValues = hyperparamValues;
info.cvAccuracy     = cvAccuracy;
info.bestHyperparam = bestMaxSplits;
info.cvTimeSeconds  = cvTimeSeconds;
info.trainTimeSeconds = trainTimeSeconds;
end

```

%-----

%TRAIN_RANDOM_FOREST Train a Random Forest (bagged trees) with simple CV

```

function [rfModel, info] = train_random_forest(XTrain, yTrain)

fprintf('\n==== Training Random Forest ====\n');

% Hyperparameter grid
numTreesGrid = [50, 100, 200];
kFolds = 5;
cvAccuracy = zeros(size(numTreesGrid));

% Time the cross validation
tCV = tic;
for i = 1:numel(numTreesGrid)
    params.numTrees = numTreesGrid(i);
    fprintf(' CV for NumTrees = %d ... \n', params.numTrees); cvAccuracy(i) =
        cross_validation(XTrain, yTrain, 'rf', params, kFolds);
end
cvTimeSeconds = toc(tCV);

```

```

% Choose best hyperparameter
[~, bestIdx] = max(cvAccuracy);
bestNumTrees = numTreesGrid(bestIdx);
fprintf('Best NumTrees = %d (CV accuracy = %.3f)\n', ...
    bestNumTrees, cvAccuracy(bestIdx));

% Time the final training
tTrain = tic;
rfModel = TreeBagger(bestNumTrees, XTrain, yTrain, ...
    'Method', 'classification', ...
    'OOBPrediction', 'On');
trainTimeSeconds = toc(tTrain);

% Fill info structure
info.modelName      = 'Random Forest';
info.hyperparamName = 'NumTrees';
info.hyperparamValues = numTreesGrid;
info.cvAccuracy     = cvAccuracy;
info.bestHyperparam = bestNumTrees;
info.cvTimeSeconds  = cvTimeSeconds;
info.trainTimeSeconds = trainTimeSeconds;
end

```

%-----

```

%EVALUATE_MODEL Evaluate model on train and test sets, save confusion matrix.
function results = evaluate_model(model, XTrain, yTrain, XTest, yTest, modelName, confusionFigPath)

fprintf('\n==== Evaluating %s ===\n', modelName);

% Predictions on train data set
yPredTrain = local_predict(model, XTrain);
trainMetrics = metrics(yTrain, yPredTrain);
fprintf('%s - Train accuracy: %.4f\n', modelName, trainMetrics.accuracy);

% Predictions on test data set
yPredTest = local_predict(model, XTest);
testMetrics = metrics(yTest, yPredTest);
fprintf('%s - Test accuracy: %.4f\n', modelName, testMetrics.accuracy);

% Plot confusion matrix for test set
figure('Visible','off');
confusionchart(testMetrics.confusion, string(testMetrics.classes));
title(sprintf('Confusion Matrix - %s (Test)', modelName));

```

```

if ~isempty(confusionFigPath)
    [figDir, ~, ~] = fileparts(confusionFigPath); if
        ~exist(figDir, 'dir'); mkdir(figDir); end
    saveas(gcf, confusionFigPath);
end
close(gcf);

```

% Package results

```

results = struct();
results.modelName = modelName;
results.trainAcc = trainMetrics.accuracy;
results.testAcc = testMetrics.accuracy;
results.metricsTest = testMetrics;
results.metricsTrain = trainMetrics;
end

```

% -----

% Local helper for predictions

```

function yPred = local_predict(model, X) if
    isa(model, 'TreeBagger')
        yPred = predict(model, X);
        if iscell(yPred)
            yPred = str2double(yPred);
        end
    else
        % Assume it's a ClassificationTree or similar
        yPred = predict(model, X);
    end
end

```

%-----

%PLOT_RESULTS Plot accuracy comparison and hyperparameter tuning curves.

```
function plot_results(dtResults, rfResults, dtInfo, rfInfo, accFigPath, hyperFigPath)
```

% Accuracy comparison

```

figure('Visible','off');
modelNames = {dtResults.modelName, rfResults.modelName};
testAcc = [dtResults.testAcc, rfResults.testAcc];

bar(testAcc);
set(gca, 'XTickLabel', modelNames);
ylabel('Test accuracy');

```

```

title('Decision Tree vs Random Forest on MNIST');

for i = 1:numel(testAcc)
    text(i, testAcc(i), sprintf('%.3f', testAcc(i)), ...
        'HorizontalAlignment','center', 'VerticalAlignment','bottom');
end

if ~isempty(accFigPath)
    [figDir, ~, ~] = fileparts(accFigPath);
    if ~exist(figDir, 'dir'); mkdir(figDir); end
    saveas(gcf, accFigPath);
end
close(gcf);

% Hyperparameter curves
figure('Visible','off');
hold on;

plot(dtInfo.hyperparamValues, dtInfo.cvAccuracy, '-o', 'DisplayName', dtInfo.modelName);
plot(rfInfo.hyperparamValues, rfInfo.cvAccuracy, '-s', 'DisplayName', rfInfo.modelName);

xlabel('Hyperparameter value');
ylabel('CV accuracy');
title('Hyperparameter Tuning');
legend('Location','best');
grid on;

if ~isempty(hyperFigPath)
    [figDir, ~, ~] = fileparts(hyperFigPath); if
        ~exist(figDir, 'dir'); mkdir(figDir); end
    saveas(gcf, hyperFigPath);
end
close(gcf);
end

%-----



% IN3300 Project – Decision Tree vs Random Forest on MNIST
clear; clc; close all;
rng(42);

% filepath management  thisFile
= mfilename('fullpath');
[thisDir,~,~] = fileparts(thisFile);

```

```

addpath(fullfile(thisDir, 'utils'));

% load_data.m
[XTrain, yTrain, XTest, yTest] = load_data();

% preprocess_data.m
[XTrainProc, XTestProc] = preprocess_data(XTrain, XTest);

%train_decision_tree.m%
train_random_forest.m
[dtModel, dtInfo] = train_decision_tree(XTrainProc, yTrain);
[rfModel, rfInfo] = train_random_forest(XTrainProc, yTrain);

% time for training and CV
fprintf('\n==== Training times ====\n');
fprintf('Decision Tree - train: %.3f s, CV: %.3f s\n', ...
    dtInfo.trainTimeSeconds, dtInfo.cvTimeSeconds);
fprintf('Random Forest - train: %.3f s, CV: %.3f s\n', ...
    rfInfo.trainTimeSeconds, rfInfo.cvTimeSeconds);

% print hyperparameter tuning results
fprintf('\n==== Hyperparameter Tuning Results ====\n');

fprintf('\nDecision Tree (MaxNumSplits):\n');
for i = 1:numel(dtInfo.hyperparamValues)
    fprintf(' MaxNumSplits = %d: CV Accuracy = %.4f\n', ...
        dtInfo.hyperparamValues(i), dtInfo.cvAccuracy(i));
end

fprintf('\nRandom Forest (NumTrees):\n');
for i = 1:numel(rfInfo.hyperparamValues)
    fprintf(' NumTrees = %d: CV Accuracy = %.4f\n', ...
        rfInfo.hyperparamValues(i), rfInfo.cvAccuracy(i));
end

% evaluate_model.m
dtResults = evaluate_model(dtModel, XTrainProc, yTrain, ...
    XTestProc, yTest, 'Decision Tree', ...
    fullfile('results','confusion_dt.png'));

rfResults = evaluate_model(rfModel, XTrainProc, yTrain, ...
    XTestProc, yTest, 'Random Forest', ...
    fullfile('results','confusion_rf.png')));

% plot_results.m

```

```
resultsDir = fullfile('results');
if ~exist(resultsDir, 'dir'); mkdir(resultsDir); end

plot_results(dtResults, rfResults, dtInfo, rfInfo, ...
    fullfile('results','accuracy_comparison.png'), ...
    fullfile('results','hyperparameter_plot.png'));
```